Robot Dynamics for Embodied Intelligence

Course home · 简体中文 · PDF

Item	Details
Applicability	Rigid-body theory; MuJoCo and Isaac examples are conceptual, not a tested cross-engine pair
Documentation review date	2026-10-18
License	CC BY 4.0 · Xiong Zeng and contributors

Verification scope: Documentation organization and upstream reference snapshot; no tutorial runtime, training or hardware validation. See version notes for pinned source references. Original teaching content uses CC BY 4.0; third-party quotations, figures and code retain their original licenses. Citation alone does not grant relicensing rights.

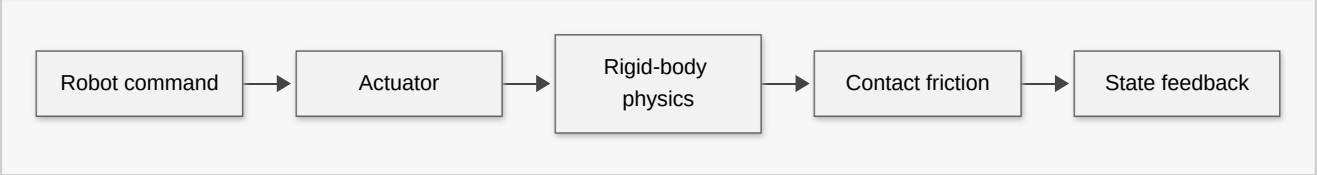
Contents

- Reading guide: follow the physical chain, not a collection of formulas
- Chapter 1. Dynamics through an embodied manipulation task
  - 1.1 Kinematics says where; dynamics explains motion
  - 1.2 State, action, and observation are different quantities
  - 1.3 One equation organizes the entire tutorial
- Chapter 2. Only the kinematics needed for dynamics
  - 2.1 Joint coordinates versus end-effector coordinates
  - 2.2 The Jacobian connects velocity and force
  - 2.3 Why inverse kinematics cannot replace dynamics
- Chapter 3. Mass, inertia, and gravity from one rigid link
  - 3.1 Translation: mass resists acceleration
  - 3.2 Rotation: inertia depends on mass distribution
  - 3.3 A one-joint pendulum already contains the main ideas
  - 3.4 Velocity-dependent coupling in three-dimensional rotation
- Chapter 4. From one link to the manipulator equation
  - 4.1 Why a mass matrix appears
  - 4.2 Deriving the standard equation from energy
  - 4.3 Forward dynamics versus inverse dynamics
  - 4.4 Energy provides a useful consistency check
- Chapter 5. A complete two-link planar manipulator example
  - 5.1 Geometry and sign conventions
  - 5.2 Kinetic energy reveals joint coupling
  - 5.3 Velocity coupling and gravity

    - 5.4 Four thought experiments worth doing
- Chapter 6. How dynamics scales to complex robots
    - 6.1 Why we do not expand every rigid-body equation
    - 6.2 Gravity compensation is a special inverse-dynamics case
- Chapter 7. Actuators: how policy actions reach the real robot
    - 7.1 A position target is not an instantaneous position jump
    - 7.2 Nonideal effects that materially change behavior
    - 7.3 Map policy action spaces to physical interfaces
- Chapter 8. Contact dynamics: grasping, pushing, sliding, and dropping
    - 8.1 Where does contact force come from?
    - 8.2 Why contact is unilateral
    - 8.3 The friction cone determines whether an object stays grasped
    - 8.4 An impact produces a rapid velocity change
    - 8.5 Distinguish four modes during a grasp
- Chapter 9. Dexterous manipulation is a coupled arm–hand–object system
    - 9.1 Why controlling the fingers alone may be insufficient
    - 9.2 When an object may be treated as extra payload
    - 9.3 Practical lessons for Tron2 with a Revo3 hand
- Chapter 10. Humanoids: from fixed-base to floating-base dynamics
    - 10.1 Why the torso cannot be driven like an ordinary joint
    - 10.2 Balance requires external support forces
    - 10.3 Only the essential centroidal dynamics and WBC concepts
- Chapter 11. Dynamics-based control that matters for embodied intelligence
    - 11.1 PD with and without gravity compensation
    - 11.2 Computed-torque control
    - 11.3 Impedance control supports compliant contact
    - 11.4 Where MPC and WBC fit
- Chapter 12. MuJoCo versus Isaac Sim / Isaac Lab: appropriate roles
    - 12.1 What each platform contributes
    - 12.2 Matching model filenames is not enough
    - 12.3 Three time scales must be identified separately
- Chapter 13. Code-free laboratory guide: free-space mechanics first
    - Lab 1 | One-link gravity and inertia
    - Lab 2 | Two-link mass matrix and coupling
    - Lab 3 | Forward–inverse dynamics consistency
    - Lab 4 | PD control versus gravity-compensated PD
    - Lab 5 | Action rate and actuator saturation
- Chapter 14. Code-free laboratory guide: contact, hands, and humanoids
    - Lab 6 | A falling block: normal contact and impact
    - Lab 7 | Sliding on a slope: friction intuition

    - Lab 8 | Dexterous grasp, rotation, and drop recovery
    - Lab 9 | Humanoid standing: contact force and center of mass
    - Lab 10 | Cross-simulator and sim-to-real sensitivity
    - Minimal one-page lab-report structure
- Chapter 15. From simulation to real robots: a concise checklist
    - 15.1 Physical parameters must remain self-consistent
    - 15.2 Prioritize sim-to-real variation by plausible error sources
    - 15.3 Four misconceptions to avoid
    - 15.4 One-page equation reference
- Appendix A. What was kept from the original course
- Appendix B. Notation, reading route, and references

Chapter concept map



Reading guide: follow the physical chain, not a collection of formulas

Robot dynamics answers an immediate question in embodied intelligence: **after a policy issues an action, why do the robot and its environment move as they do?** A single task connects the entire tutorial: a robot arm with a dexterous hand picks up a screwdriver, rotates it, and picks it up again after a drop. A humanoid that stands and reaches for the same object follows the same physical laws, except that its base is floating and its feet establish additional contacts.

The tutorial draws on the structure of Wei Zhang's MEE5114 Advanced Control for Robotics, especially the lectures on rigid-body motion, dynamics, and control. It is **not a reproduction or translation of those lecture notes**. We also use *Modern Robotics*, Featherstone's rigid-body dynamics work, and official MuJoCo / Isaac Lab documentation, focusing only on material with a clear role in robot learning.[1–8]

The central chain is: policy action a → actuator torque τ → acceleration $\ddot{q}$ and contact force λ → next physical state → camera and proprioceptive observation o → next policy action.

$$\boxed{a_t \longrightarrow \tau_t \longrightarrow (\ddot{q}_t, \lambda_t) \longrightarrow s_{t+1} \longrightarrow o_{t+1}.}$$

Included: the minimum coordinate and Jacobian concepts; mass, inertia, and gravity; manipulator equations; forward and inverse dynamics; contact and friction; actuation; floating-base dynamics; control; simulation; and sim-to-real transfer.

Deliberately omitted: Lie-group proofs, full screw theory and product-of-exponentials derivations, detailed spatial cross-product operators, low-level implementations of RNEA/ABA, general-purpose

SDP and optimal-control theory, software installation commands, and all programming code. Readers who need those topics can proceed to the references.

How to read: Chapters 1–6 build the physics foundation. Chapters 7–11 explain actuation, contact, and control. Chapters 12–14 specify reproducible **code-free** MuJoCo and Isaac Sim / Isaac Lab experiments. Chapter 15 consolidates sim-to-real considerations. The experiments describe what to perform and measure; they are not claims that simulations have already been executed.

Chapter 1. Dynamics through an embodied manipulation task

1.1 Kinematics says where; dynamics explains motion

Let $q \in \mathbb{R}^n$ denote the joint configuration and $\dot{q}$ the joint velocity. Kinematics asks: where is the palm given q, and how fast does it move given $\dot{q}$? It requires no mass, inertia, or motor-torque information.

Dynamics asks: what acceleration follows from an applied torque? How does the table's support force alter the finger motion? What happens after the object leaves contact? Dynamics thus concerns the **state, input forces, and environmental interaction**, rather than geometry alone.

Question in the task	Main concept	Why it matters
Where should a fingertip go?	Forward/inverse kinematics	Convert object pose into target joint motion
What actually drives the finger?	Actuators + dynamics	Explain torque, acceleration, and delay
Why does the screwdriver stay in the hand?	Contacts + friction	Check whether forces lie inside the friction cone
Why does a replacement hand change tracking?	Mass, inertia, gravity	Payload and center of mass have changed
Why do humanoid feet need ground forces?	Floating base + centroidal dynamics	Internal motors cannot suspend the whole body

1.2 State, action, and observation are different quantities

A physical state may include robot joint positions and velocities, object pose and velocity, and contact mode: $s = (q, \dot{q}, x_o, v_o, \text{contacts}, \ldots)$. A learned policy receives observations o, such as RGB-D frames, joint encoders, and estimated finger currents. A camera does not directly measure the mass matrix or the true friction coefficient.

The meaning of an action a depends on the control interface. It may specify desired joint position, velocity, torque, or an end-effector pose. **The same numerical action can yield different physical motions under different low-level servos.** This is the first question to ask when comparing robot-learning action spaces.

1.3 One equation organizes the entire tutorial

For a fixed-base, fully actuated rigid manipulator with independent joint coordinates,

$$\boxed{M(q)\ddot{q} + C(q,\dot{q})\dot{q} + g(q) = \tau + J_c(q)^T\lambda.}$$

Here M is configuration-dependent inertia, $C\dot{q}$ captures velocity-dependent coupling, g is the generalized gravity term, τ is actuator torque, λ is the force of the environment on the robot, and $J_c^T\lambda$ maps that contact force into generalized coordinates.

Define $h(q,\dot{q}) = C(q,\dot{q})\dot{q} + g(q)$. The compact form is

$$M(q)\ddot{q} + h(q,\dot{q}) = \tau + J_c^T\lambda.$$

Real systems may also contain friction, compliant transmissions, saturation, and delay. Those effects must be restored when interpreting experimental discrepancies.

Chapter 2. Only the kinematics needed for dynamics

2.1 Joint coordinates versus end-effector coordinates

For n joints, $q = (q_1, \ldots, q_n)$ describes the configuration. End-effector position is $p = f(q)$. To represent a full six-degree-of-freedom pose, use

$$T = \begin{bmatrix} R & p \\ 0 & 1 \end{bmatrix}, \qquad R^TR = I,\ \det R = 1.$$

R is orientation; p is translation. Frame labels matter: T_{AB} maps coordinates expressed in frame B into frame A. For embodied learning, **consistent camera, robot, object, and table frames** are more important than proving every property of $SO(3)$ or $SE(3)$.

2.2 The Jacobian connects velocity and force

For end-effector position $p(q)$,

$$\boxed{\dot{p} = J_p(q)\dot{q}, \qquad \ddot{p} = J_p(q)\ddot{q} + \dot{J}_p(q,\dot{q})\dot{q}.}$$

The Jacobian first maps joint velocity to end-effector velocity. More importantly, an external Cartesian force f generates joint-space generalized force

$$\boxed{\tau_{\text{ext}} = J_p(q)^T f.}$$

The reason is **virtual work**: $f^T\delta p = f^T J_p \delta q = \tau_{\text{ext}}^T \delta q$. Thus, the table's force on a fingertip affects every joint through this mapping.

For full hand orientation and translation, one may use a six-dimensional rigid-body velocity (twist) and a six-dimensional Jacobian. The essential caution is to keep the **reference point, velocity representation, and wrench representation consistent**; a complete derivation of adjoint transformations is unnecessary here.

2.3 Why inverse kinematics cannot replace dynamics

Inverse kinematics converts a target end-effector pose into joint targets. It does not account automatically for object weight, rotational inertia, or contact forces. A geometrically feasible path may still be unrealizable within motor limits and time constraints. Picking up a heavy object changes the torques needed to hold the same pose.

Chapter 3. Mass, inertia, and gravity from one rigid link

3.1 Translation: mass resists acceleration

The center of mass of a rigid body satisfies Newton's second law:

$$\boxed{f = ma_C.}$$

For the same applied force, a larger mass generally produces a smaller center-of-mass acceleration. In grasping, identical policies can accelerate or drop objects differently when object mass changes.

3.2 Rotation: inertia depends on mass distribution

For rotation about a fixed axis, $\tau = I_O\ddot{q}$ when other torques are ignored. The rotational inertia I_O depends on **where the mass lies relative to the axis**, not just on total mass. If the center of mass is at distance r from the joint axis, the parallel-axis theorem gives

$$\boxed{I_O = I_C + mr^2.}$$

Two hands with identical masses but different mounting offsets can produce different effective joint inertia. Replacing an end-effector can therefore alter the dynamic behavior of the entire arm.

3.3 A one-joint pendulum already contains the main ideas

Consider a rigid link rotating in a vertical plane, with positive angle q measured upward from the horizontal. Its center of mass is distance r from the axis; its moment of inertia about that axis is I_O. Ignoring friction,

$$\boxed{I_O\ddot{q} + mgr\cos q = \tau.}$$

This includes the inertia term $I_O\ddot{q}$, gravity term $mgr\cos q$, and drive torque τ. A horizontal stationary link requires holding torque mgr; a vertically oriented link requires zero *static gravity* torque.

Key distinction: gravity compensation is configuration-dependent, not a fixed offset added to every joint. Being stationary does not imply zero motor torque.

3.4 Velocity-dependent coupling in three-dimensional rotation

For a rigid body expressed in a body-fixed frame, rotational motion satisfies

$$\tau_C = I_C\dot{\omega} + \omega \times (I_C\omega).$$

The second term reflects changes in angular momentum direction. It is enough to understand why the term exists: in embodied-learning practice, a multibody solver usually handles this coupling without requiring manual expansion of every link's three-dimensional equations.

Chapter 4. From one link to the manipulator equation

4.1 Why a mass matrix appears

Motion of one joint may carry many downstream links. Effective inertia therefore depends on the robot's entire configuration. Summing translational and rotational kinetic energies yields the quadratic form

$$\boxed{K(q,\dot{q}) = \tfrac{1}{2}\dot{q}^T M(q)\dot{q}.}$$

The generalized mass matrix $M(q)$ is generally **not diagonal**. Its off-diagonal entries express inertial coupling between joints. For independent generalized coordinates of a nondegenerate rigid mechanical system, $M(q)$ is symmetric positive definite.

4.2 Deriving the standard equation from energy

Let $U(q)$ be gravitational potential energy, and define the Lagrangian $L = K - U$. The Euler–Lagrange equations are

$$\frac{d}{dt}\frac{\partial L}{\partial \dot{q}_i} - \frac{\partial L}{\partial q_i} = \tau_i.$$

Substituting quadratic kinetic energy and collecting terms gives

$$\boxed{M(q)\ddot{q} + C(q,\dot{q})\dot{q} + g(q) = \tau, \qquad g(q) = \nabla_q U(q).}$$

The velocity term $C\dot{q}$ contains squared velocities and products of different joint velocities. It is not an arbitrary damper or an external contact force; it comes from rigid-body motion itself. In practice, the bias vector $h = C\dot{q} + g$ is often used directly; the matrix representation of C is not unique.

4.3 Forward dynamics versus inverse dynamics

Forward dynamics predicts acceleration from state and applied torque:

$$\boxed{\ddot{q} = M(q)^{-1}[\tau - h(q,\dot{q})] \quad \text{(without contacts).}}$$

Inverse dynamics finds the torque required for a desired acceleration:

$$\boxed{\tau = M(q)\ddot{q} + h(q,\dot{q}).}$$

Physics simulation primarily uses the forward problem; gravity compensation, feedforward tracking, and certain whole-body controllers use the inverse problem. The inverse matrix above is conceptual notation; numerical implementations typically solve the associated linear system rather than explicitly computing M^{-1}.

4.4 Energy provides a useful consistency check

Without friction or external contact, total mechanical energy $E = K + U$ satisfies

$$\boxed{\dot{E} = \tau^T \dot{q}.}$$

Motor mechanical power changes stored mechanical energy. Dissipative friction removes energy; moving external contacts can add or remove energy through contact work. This identity helps

explain simulation energy drift, controller behavior, and compliance, without requiring a full stability-theory treatment.

Chapter 5. A complete two-link planar manipulator example

5.1 Geometry and sign conventions

Consider a two-revolute-joint arm moving in a vertical plane. q_1 is measured upward from the horizontal; q_2 is relative to link 1. Link 1 has length l_1, mass m_1, center-of-mass distance r_1 from joint 1, and center-of-mass inertia I_1. Link 2 has corresponding parameters m_2, r_2, I_2. Its absolute orientation is $q_1 + q_2$.

The center-of-mass positions are

$$\begin{aligned} p_1 &= (r_1 \cos q_1,\ r_1 \sin q_1)^T, \\ p_2 &= (l_1 \cos q_1 + r_2 \cos(q_1 + q_2), \\ &\qquad l_1 \sin q_1 + r_2 \sin(q_1 + q_2))^T. \end{aligned}$$

I_1 and I_2 are the link inertias about their centers of mass and normal to the plane of motion.

5.2 Kinetic energy reveals joint coupling

Total kinetic energy is

$$K = \frac{1}{2} m_1 \|\dot{p}_1\|^2 + \frac{1}{2} I_1 \dot{q}_1^2 + \frac{1}{2} m_2 \|\dot{p}_2\|^2 + \frac{1}{2} I_2 (\dot{q}_1 + \dot{q}_2)^2.$$

Collecting it as $K = \frac{1}{2} \dot{q}^T M(q) \dot{q}$ gives

$$\begin{aligned} M_{11} &= I_1 + I_2 + m_1 r_1^2 + m_2 (l_1^2 + r_2^2 + 2 l_1 r_2 \cos q_2), \\ M_{12} &= M_{21} = I_2 + m_2 (r_2^2 + l_1 r_2 \cos q_2), \\ M_{22} &= I_2 + m_2 r_2^2. \end{aligned}$$

Read the physics in the matrix: M_{11} varies with q_2, and M_{12} is usually nonzero. These are configuration-dependent inertia and cross-joint inertial coupling, respectively.

5.3 Velocity coupling and gravity

Let $b = m_2 l_1 r_2$. Then

$$C(q, \dot{q}) \dot{q} = \begin{bmatrix} -b \sin q_2 (2 \dot{q}_1 \dot{q}_2 + \dot{q}_2^2) \\ b \sin q_2 \dot{q}_1^2 \end{bmatrix}.$$

Gravitational potential energy is

$$U = m_1 g r_1 \sin q_1 + m_2 g [l_1 \sin q_1 + r_2 \sin(q_1 + q_2)],$$

so

$$g(q) = \begin{bmatrix} (m_1 r_1 + m_2 l_1) g \cos q_1 + m_2 r_2 g \cos(q_1 + q_2) \\ m_2 r_2 g \cos(q_1 + q_2) \end{bmatrix}.$$

The final result is $\tau = M(q)\ddot{q} + C(q,\dot{q})\dot{q} + g(q)$. We have followed the complete chain of center-of-mass positions → velocities → kinetic and potential energy → mass matrix and gravity vector. More elaborate symbolic three-dimensional calculations are unnecessary for understanding this structure.

5.4 Four thought experiments worth doing

1. **Stationary, no contact**: $\dot{q} = \ddot{q} = 0$ requires $\tau = g(q)$.
2. **Slow motion**: velocity-dependent terms are small relative to the most important inertial and gravity effects.
3. **Fast motion**: quadratic velocity terms become significant; neglecting $C\dot{q}$ introduces prediction error.
4. **A heavier second link**: both M and g change, so the same position-controller gains may need retuning.

Chapter 6. How dynamics scales to complex robots

6.1 Why we do not expand every rigid-body equation

A robot with dozens of joints is cumbersome to analyze through full symbolic expansion. Modern dynamics software exploits the kinematic tree and uses recursive algorithms. For embodied intelligence, the most important thing is to know **what each algorithm computes**, rather than memorizing its spatial-vector implementation.

Algorithm	Inputs	Output	Use in robot learning
RNEA: recursive Newton–Euler	$q, \dot{q}, \ddot{q}$	Inverse dynamics force	Gravity compensation and torque feedforward
CRBA: composite rigid-body algorithm	q	$M(q)$	Inertial coupling and model-based control
ABA: articulated-body algorithm	$q, \dot{q}, \tau$	$\ddot{q}$	Forward dynamics and physics simulation

RNEA propagates velocity and acceleration from base to tip, then accumulates forces back toward the base. ABA uses articulated-body inertia to solve accelerations efficiently. CRBA assembles the composite inertia of downstream links. With matching models and conventions in free space, forward and inverse results should agree.

6.2 Gravity compensation is a special inverse-dynamics case

Setting desired joint velocity and acceleration to zero gives, for an ideal rigid manipulator,

$$\boxed{g(q) = \mathrm{RNEA}(q, 0, 0).}$$

This is physically more appropriate than setting a constant torque at every joint. If the robot carries a payload, its mass, center of mass, and inertia must be included in the model; an algorithm can be numerically correct yet predict the wrong physics when its parameters are wrong.

Chapter 7. Actuators: how policy actions reach the real robot

7.1 A position target is not an instantaneous position jump

A robot-learning policy often outputs q_{target}. The motor does not instantly transport the robot to that configuration. A representative joint PD actuator uses

$$\boxed{\tau_{\text{servo}} = K_p(q_{\text{target}} - q) + K_d(\dot{q}_{\text{target}} - \dot{q}).}$$

The physical dynamics then determine $\ddot{q}$. The actual chain is target position → servo torque → acceleration → velocity → position. Some real robots hide this servo within a drive, exposing only a position interface to the learning system.

7.2 Nonideal effects that materially change behavior

Effect	Physical consequence	Training/deployment failure
Torque or current saturation	Acceleration cannot grow indefinitely	Requested actions are not realized
Velocity limit	Very fast motion is clipped	Tracking lag and contact timing shifts
Friction and backlash	Small commands may produce no motion	Poor precision at low speed
Delay and filtering	Commands act later than expected	Overshoot and oscillation
Motor and payload inertia	Dynamic response changes with the tool	Same gains, different results
Compliant transmission	Joints and fingers deform elastically	Different contact-force response

For an ideal gear ratio N, motor rotor inertia I_m reflects approximately as $N^2 I_m$ at the joint. Link inertia alone thus cannot characterize actuator dynamics. Real drives may also include current, velocity, and position loops plus thermal protection.

7.3 Map policy action spaces to physical interfaces

Policy action	Advantage	Physical caveat
Target joint position	Stable and common in imitation learning	Depends on servo gains and command period
Joint-position increment	Convenient for residual tracking	Drift and joint-limit handling
Target joint velocity	Directly expresses motion direction	Integration and low-level-drive differences
Joint torque	Direct interaction with dynamics	Harder training, saturation and safety
End-effector pose increment	Task-relevant representation	Needs inverse kinematics and contact handling

Conclusion: when comparing robot-learning algorithms, report the low-level actuator model as well as the nominal action space; otherwise the same action symbol may represent different physical control systems.

Chapter 8. Contact dynamics: grasping, pushing, sliding, and dropping

8.1 Where does contact force come from?

If a fingertip remains in fixed contact with a table, the contact point has zero velocity relative to the table:

$$J_c(q)\dot{q} = 0, \qquad J_c(q)\ddot{q} + \dot{J}_c(q,\dot{q})\dot{q} = 0.$$

Dynamics must also include the generalized contact force:

$$\boxed{M(q)\ddot{q} + h(q,\dot{q}) = \tau + J_c(q)^T\lambda.}$$

Acceleration and contact force are therefore **coupled**. If the robot tries to move into the table, the table supplies a normal force that prevents penetration and changes the resulting motion and joint torques. The zero-relative-velocity constraint applies only to a sustained sticking contact, not to every contact mode.

8.2 Why contact is unilateral

Let the normal gap between surfaces be $\phi(q) \geq 0$ and the compressive normal force be $\lambda_n \geq 0$. Ideal rigid unilateral contact satisfies

$$\boxed{\phi(q)\lambda_n = 0.}$$

This says: **either the surfaces are separated and the force is zero, or they touch and can sustain compression**. A normal table cannot keep pulling a screwdriver downward after it has lifted away. Adhesion, suction, and magnetism require separate models.

8.3 The friction cone determines whether an object stays grasped

For Coulomb friction, tangential contact force λ_t and normal force λ_n satisfy

$$\boxed{\|\lambda_t\| \leq \mu\lambda_n.}$$

μ is the friction coefficient. Increasing normal grip force generally increases the maximum available static friction, subject to motor limits, geometry, deformation, and safety. An object can slip because of low normal force, low friction, changing gravity direction, or increased inertial load during rapid manipulation.

8.4 An impact produces a rapid velocity change

When a finger collides with an object, large forces may act over a short interval. An impulse P approximates the change in generalized velocity across impact:

$$\boxed{M(q)(v^+ - v^-) = J_c(q)^T P.}$$

Contact solvers handle not only sustained support forces but also velocity changes during collisions. Different simulators may employ different contact compliance, time discretization, friction

approximations, and constraint solvers. Equal masses and friction coefficients do not guarantee identical collision histories.

8.5 Distinguish four modes during a grasp

Mode	Observable behavior	Implication for the policy
No contact	A gap exists between finger and object	Approach before exerting grasp force
Sticking contact	Contact points have near-zero relative velocity	A stable manipulation constraint can form
Sliding contact	Contact points move tangentially	Adjust grip, trajectory, or planning
Separation/drop	Contact forces disappear; object moves independently	Relocalize, approach, and regrasp

This explains why **drop recovery is not ordinary trajectory tracking**: the contact mode has changed, and the controller must establish new contacts rather than continue an obsolete motion reference.

Chapter 9. Dexterous manipulation is a coupled arm–hand–object system

9.1 Why controlling the fingers alone may be insufficient

Let q_r be generalized robot coordinates, including arm and hand, and let x_o be the object pose. Fingertip forces act on both the robot and the object, with action and reaction. On the robot side,

$$M_r(q_r)\ddot{q}_r + h_r = \tau_r + J_r^T \lambda.$$

On the object side, linear and angular momentum change in response to gravity and fingertip contact forces. Force signs must be chosen consistently: the two bodies do not both experience the same contact force in the same direction.

When the palm transports a rigidly held object, combined payload mass, center of mass, and inertia dominate the arm response. When the object rolls or rotates within the fingers, contact locations, friction, and object orientation become central. The arm, hand, and object should be treated as **one coupled physical system**.

9.2 When an object may be treated as extra payload

If an object remains firmly grasped and scarcely moves relative to the palm, it can be approximated as a rigid additional payload. Updating the total end-effector mass, center of mass, and inertia then supports gravity compensation and slow transport.

The approximation fails when the object rotates in the fingers, slides, changes contact points, or drops. Those phases require explicit object degrees of freedom and contact constraints.

9.3 Practical lessons for Tron2 with a Revo3 hand

After replacing a dexterous hand, verify at least: mounting frame and tool center point (TCP); payload mass, center of mass, and inertia; joint gains and force limits; fingertip collision geometry; fingertip–object friction; observation rate and command delay.

Changing the TCP is primarily a **geometric and Jacobian change**, not a change in physical inertia. Replacing the actual hand changes the physical payload. Head and wrist cameras provide object-pose observations but do not directly measure accurate contact forces. Without tactile sensors, finger current, joint behavior, and visual motion can provide uncertain proxies for contact.

Chapter 10. Humanoids: from fixed-base to floating-base dynamics

10.1 Why the torso cannot be driven like an ordinary joint

A table-mounted robot arm has a fixed base. A humanoid torso can translate and rotate freely in space, yet it generally has no motors that directly actuate these six world-space base degrees of freedom. For n joints, floating-base generalized velocity has dimension $n+6$. If orientation is represented by a quaternion, configuration variables typically have dimension $n+7$; configuration and velocity vectors are not interchangeable.

The dynamics become

$$\boxed{M(q)\dot{v} + h(q,v) = S^T\tau + J_c(q)^T\lambda,}$$

where S inserts actuated joint torques into the full generalized force vector. Base motion depends on the coupled effects of joints and external contacts.

10.2 Balance requires external support forces

For the whole-robot center of mass c,

$$\boxed{m\ddot{c} = m\mathbf{g} + \sum_i f_i.}$$

During quiet standing, $\ddot{c} \approx 0$, so ground contact forces approximately balance gravity. Without foot contact, internal joint torques cannot prevent the entire robot's center of mass from falling. Reaching, walking, and dexterous grasping all involve **the relationship between external force and available actuation**.

10.3 Only the essential centroidal dynamics and WBC concepts

The change in angular momentum about the center of mass follows external moments:

$$\dot{L}_c = \sum_i (r_i - c) \times f_i + \sum_i m_i^{\text{ext}}.$$

Whole-body control (WBC) considers end-effector tracking, torso balance, fixed foot contacts, actuator limits, friction cones, and joint limits together. Model-based WBC can explicitly optimize accelerations and contact forces; an RL locomotion policy can instead learn behavior through

simulation and reward. **Both rely on the same underlying physics**, but a learned policy need not solve inverse dynamics or a quadratic program during deployment.

Chapter 11. Dynamics-based control that matters for embodied intelligence

11.1 PD with and without gravity compensation

Given desired joint position q_d and velocity $\dot{q}_d$, a standard torque law is

$$\tau = K_p(q_d - q) + K_d(\dot{q}_d - \dot{q}).$$

Adding model-based gravity compensation gives

$$\boxed{\tau = K_p(q_d - q) + K_d(\dot{q}_d - \dot{q}) + g(q).}$$

This is the simplest way to understand a position-action interface's low-level feedback loop. It also explains why changing payloads often requires model or gain updates. Gravity compensation alone does not guarantee stable contact behavior.

11.2 Computed-torque control

In free space, with accurate dynamics and accessible torque control, let $e = q_d - q$ and define

$$\ddot{q}^* = \ddot{q}_d + K_d\dot{e} + K_p e.$$

Choose

$$\boxed{\tau = M(q)\ddot{q}^* + h(q, \dot{q}).}$$

Under ideal assumptions the resulting error dynamics are $\ddot{e} + K_d\dot{e} + K_p e = 0$. The derivation shows how a model can compensate nonlinear mechanical coupling to yield simple tracking-error dynamics. Real robots remain subject to saturation, friction, inaccurate parameters, and contact disturbances.

11.3 Impedance control supports compliant contact

In dexterous manipulation, infinitely stiff position tracking is undesirable. A simplified desired Cartesian impedance force is

$$F_d = K_x(x_d - x) + D_x(\dot{x}_d - \dot{x}),$$

which is mapped to joint torque through the Jacobian and combined with gravity or other dynamic terms as necessary. Compared with rigid tracking, impedance control deliberately permits some displacement in response to contact force, supporting safer and more compliant interaction.

11.4 Where MPC and WBC fit

Model predictive control (MPC) predicts several future steps using dynamics and optimizes control inputs; WBC concentrates on satisfying full-body motion tasks and contact constraints at once. A

detailed optimization-solver derivation is beyond this tutorial. The important principle is that **physical constraints determine which desired accelerations, torques, and contact forces can coexist**. High-level VLA action targets are not automatically dynamically feasible.

Chapter 12. MuJoCo versus Isaac Sim / Isaac Lab: appropriate roles

12.1 What each platform contributes

MuJoCo is well suited to building mechanical intuition, measuring sensitivity to physical parameters, investigating contact behavior, and running comparatively lightweight arm and hand experiments. Its efficient simulation does not eliminate the need to understand contact models and solver settings.

Isaac Sim provides robot and scene simulation capabilities. **Isaac Lab** organizes robot-learning environments, parallel experiments, action/observation/reward pipelines, and domain randomization on top of simulation capabilities. Isaac Sim and Isaac Lab are not two independent physics engines. The active physics backend, version, and actuator semantics must be taken from the documentation matching the local installation.[4–7]

Research question	MuJoCo emphasis	Isaac Sim / Isaac Lab emphasis
Why does a link fall?	Minimal model with interpretable torque terms	Verify body properties and joint drives
Why does a grasp slip?	Geometry, friction, and numerical contact settings	Multi-contact scenes, batched environments, randomization
Why does real-robot performance degrade?	Small, diagnostic parameter sweeps	Large-scale training and sim-to-real variation
How should a controller be evaluated?	Transparent single-environment baseline	Scaling to full robots and tasks

12.2 Matching model filenames is not enough

A fair cross-platform comparison must align body mass and center of mass, inertia tensor and inertial frame, joint axes and limits, base fixation, collision geometry, gravity, actuator gains and limits, physics time step, control rate, initial conditions, and contact/friction settings.

Even after careful matching, contact compliance, discretization, and constraint solvers may differ. **Compare contact-free motion first, then gradually add contact.** Neither simulator should be treated as ground truth for the other.

12.3 Three time scales must be identified separately

The physics integration step Δt_{phys}, servo/controller update period Δt_{ctrl}, and policy action period Δt_{π} may be different. A low-frequency policy may send position targets while the low-level PD loop runs multiple times between actions and the solver performs even smaller physics steps.

A fair policy comparison must report all three frequencies and how actions are held or interpolated; otherwise nominally identical policies can experience different contact transients and delays.

Chapter 13. Code-free laboratory guide: free-space mechanics first

Use **the same one-link pendulum and the same two-link arm** in both platforms as the common starting point. Each lab gives an objective, setup, measurements, and a prediction. No commands, software snippets, or executable scripts are required to understand the exercise.

Lab 1 | One-link gravity and inertia

Objective: separate the effects of mass, center of mass, and rotational inertia. **Setup:** create a single link on a fixed rotational axis in a vertical plane. Release it from a horizontal pose, then independently vary mass, center of mass, and inertia within physically plausible ranges. **Measure:** $q(t)$, $\dot{q}(t)$, $\ddot{q}(t)$, and gravitational joint torque. **Expected:** changing center of mass changes both gravity torque and joint inertia; a stationary horizontal link requires nonzero holding torque.

Check: align axis direction, gravity, and zero-angle convention in MuJoCo and Isaac Sim, or the gravity-torque curves may differ by a sign or phase.

Lab 2 | Two-link mass matrix and coupling

Objective: understand why $M(q)$ depends on configuration. **Setup:** use the two-link parameters from Chapter 5, hold q_1 fixed, and gradually vary q_2 while examining equivalent joint inertia and static holding torque. **Measure:** diagonal and off-diagonal mass-matrix terms and $g(q)$. **Expected:** M_{11} and M_{12} vary with $\cos q_2$; joints are inertially coupled.

Check: remove contact and unintended drives when comparing mass-matrix properties, and verify SI units.

Lab 3 | Forward–inverse dynamics consistency

Objective: verify $M\ddot{q} + h = \tau$. **Setup:** choose matching nonsingular robot configurations, velocities, and desired accelerations. Determine the generalized forces required for the desired acceleration, then predict free-space acceleration under those forces. **Measure:** acceleration discrepancy and equation residual. **Expected:** values are close with consistent conventions; discretization and actuator modeling introduce differences.

Check: a position-servo target is not a joint torque. Identify whether additional low-level controllers are active.

Lab 4 | PD control versus gravity-compensated PD

Objective: connect feedback to model feedforward. **Setup:** hold a two-link arm in horizontal and downward-pointing configurations; compare pure PD against PD with gravity compensation. **Measure:** steady-state error, maximum joint effort, and settling time. **Expected:** with suitable gains and a correct model, gravity compensation reduces pose-dependent error; it cannot remove all delay or unmodeled friction.

Lab 5 | Action rate and actuator saturation

Objective: see why identical policy targets can result in different motion. **Setup:** hold the target trajectory constant but vary physics step, servo rate, policy rate, torque limit, and PD gains. **Measure:**

tracking error, peak torque, overshoot, and lag. **Expected:** slower command updates or stricter limits may increase tracking error; excessively high gains can destabilize contact interactions or coarse discretization.

Chapter 14. Code-free laboratory guide: contact, hands, and humanoids

Lab 6 | A falling block: normal contact and impact

Objective: observe the transition from free fall to supported contact. **Setup:** drop blocks with the same mass and collision geometry from the same height onto a table in MuJoCo and Isaac Sim, then vary contact settings in a controlled manner. **Measure:** height, vertical speed, penetration, normal force/impulse, and resting jitter. **Expected:** impact details can differ across platforms; after settling, average normal support force should approximately equal object weight.

Lab 7 | Sliding on a slope: friction intuition

Objective: study friction coefficient and sliding threshold. **Setup:** place a block on a slowly increasing incline and repeat with different static-friction parameters. **Measure:** onset angle, sliding velocity, and normal force. **Expected:** ideal Coulomb static friction gives $\tan\theta_c \approx \mu$; static-friction approximations and time discretization can shift this threshold in real solvers.

Lab 8 | Dexterous grasp, rotation, and drop recovery

Objective: combine mass, inertia, friction, and contact-mode changes in one task. **Setup:** keep screwdriver geometry fixed, then vary object mass/center of mass, friction, initial orientation, grip location, and manipulation speed. Define phases for grasp → lift → rotate → disturb → regrasp when necessary. **Measure:** success rate, object orientation error, slip distance, contact duration, post-drop recovery rate, and number of actions.

Expected: rapid rotation raises inertial loads and low friction increases slip risk. Successful recovery does not prove that the initial grasp was optimal; check whether the policy actually detected the new contact state. Before sim-to-real on Tron2 + Revo3, verify mounting geometry, payload data, and fingertip materials.

Lab 9 | Humanoid standing: contact force and center of mass

Objective: verify that the floating base cannot be suspended by internal actuation. **Setup:** choose a humanoid or quadruped with foot contact, and compare stable standing, a small forward body lean, and lifting one foot. **Measure:** center-of-mass projection, ground reaction force, torso pose, and joint effort. **Expected:** support forces redistribute and stability requires joint behavior consistent with contact and center-of-mass mechanics.

Lab 10 | Cross-simulator and sim-to-real sensitivity

Objective: identify the physical parameters that matter most for learning. **Setup:** first align the contact-free two-link models, then introduce friction, payload, joint delay, and contact. Vary one parameter family at a time before joint domain randomization. **Measure:** dynamics residual,

trajectory error, contact duration, task success, and robustness curves. **Expected:** free-space comparisons are easier to make consistent; contact and actuator behavior are often the most difficult factors to match.

Minimal one-page lab-report structure

Record six items for every experiment: **research question, physical state and action definition, physical parameters, intervention, measured metrics, and conclusion/limitations**. At minimum, plot position and velocity over time; for contact labs add normal force or impulse. Cross-platform plots need matching units, time axes, initial conditions, and event labels.

Chapter 15. From simulation to real robots: a concise checklist

15.1 Physical parameters must remain self-consistent

Do not randomize every inertial parameter independently without constraints. Mass, center of mass, and rotational inertia are physically related. Moving a payload's mass distribution generally changes inertia as well. Inertia tensors should be symmetric and physically realizable. Collision geometry, inertial frames, and visual meshes should have clearly distinguished roles.

15.2 Prioritize sim-to-real variation by plausible error sources

For manipulation, a practical order is: payload mass and center of mass, joint friction, actuator gains and saturation, command delay, fingertip–object friction, contact-geometry error, and initial object pose. Domain randomization should cover a **plausible real-world distribution**, not simply maximize parameter ranges.

With RGB-D observations, also assess visual noise, calibration error, and occlusion separately. Perception error changes the policy's estimate of the state; dynamics error changes how the real state evolves after an action. Both matter, but they should not be conflated.

15.3 Four misconceptions to avoid

- **"The URDF is correct, so the dynamics must be correct."** No. Collision, actuator, friction, and solver properties also matter.
- **"A position-output policy does not need dynamics."** No. Position targets still act through motor forces.
- **"Gravity compensation guarantees a stable grasp."** No. Friction, payload changes, and dynamic disturbances remain.
- **"Different simulator results mean one simulator is wrong."** Not necessarily. Contact models and time-stepping methods can produce different effective behavior.

15.4 One-page equation reference

Scenario	Core relation	Interpretation
End-effector motion	$\dot{p} = J_p \dot{q}$	Joint velocity determines Cartesian velocity

Scenario	Core relation	Interpretation
End-effector force	$\tau_{\text{ext}} = J_p^T f$	Map Cartesian force to joint generalized force
Free-space manipulator	$M\ddot{q} + h = \tau$	Actuation determines joint acceleration
Manipulator with contact	$M\ddot{q} + h = \tau + J_c^T \lambda$	Contact force changes the motion
Static gravity compensation	$\tau = g(q)$ (without contact)	Holding torque depends on configuration
Friction cone	$\lvert\lambda_t\rvert \le \mu\lambda_n$	Limits sticking-contact force
Floating base	$M\dot{v} + h = S^T\tau + J_c^T\lambda$	The base is underactuated and needs external forces
Center-of-mass translation	$m\ddot{c} = m\mathbf{g} + \sum f_i$	Support forces govern standing and lifting

Appendix A. What was kept from the original course

MEE5114 material	Location here	Editorial treatment
Lectures 2–7: poses, twists, POE, Jacobians	Chapter 2	Keep frames and Jacobians; omit detailed POE derivation
Lecture 8: rigid-body dynamics	Chapters 3–4	Retain inertia and physical intuition; omit spatial-operator details
Lecture 9: dynamics of open chains	Chapters 4–6	Retain manipulator equation and roles of RNEA / CRBA / ABA
Lecture 14: robot motion control	Chapters 7 and 11	Retain servos, gravity compensation, computed torque, impedance
Added: contact and manipulation	Chapters 8–9	Essential to embodied tasks
Added: floating-base dynamics	Chapter 10	For humanoids and legged robots
Added: two-platform experiments and sim-to-real	Chapters 12–15	Objectives and measurements, with no code

Appendix B. Notation, reading route, and references

Notation: $q, \dot{q}, \ddot{q}$: joint position, velocity, acceleration; M: joint-space mass matrix; $C\dot{q}$: Coriolis/centrifugal term; $g(q)$: generalized gravity force; $h = C\dot{q} + g$: nonlinear bias; τ: actuated generalized force; J: Jacobian; λ: environmental contact force; μ: friction coefficient; v: floating-base generalized velocity; S: actuation-selection matrix.

Shortest learning path: study Chapters 1–5 and carry out Labs 1–4; continue to Chapters 7–9 with Labs 5–8. If you work on humanoids or quadrupeds, add Chapter 10 and Lab 9. Lab 10 is designed for cross-simulator consistency and sim-to-real sensitivity.

1. MEE5114 — Advanced Control for Robotics, Wei Zhang: course structure, especially Lectures 8, 9, and 14.
2. Modern Robotics — Chapter 8: Dynamics of Open Chains: Lagrangian mechanics, forward/inverse dynamics, constraints, and actuation.

3. Modern Robotics — Chapter 11: Robot Control: joint-space and task-space control.
4. Featherstone — Rigid Body Dynamics Algorithms / Spatial Vector Resources: deeper analysis of RNEA, CRBA, and ABA.
5. MuJoCo — Computation: multibody dynamics, actuators, contacts, and constraint solving.
6. Isaac Lab — Articulations: multibody structures, fixed and floating bases, action interfaces.
7. Isaac Lab — Actuators: explicit/implicit actuator semantics; match the installed version.
8. MIT — Underactuated Robotics: Multibody Dynamics: further reading on multibody and underactuated systems.